

# dk

September 26, 2017

## 1 RNN modeling of behavior and performance

Yifat Amir, Allen Guo, Sam Lau

```
In [1]: import pandas as pd
import numpy as np
import json

filename = 'skill_builder_data_corrected.csv'
df = pd.read_csv(filename, encoding='ISO-8859-1', low_memory=False)
df = df[(df['original'] == 1) & (df['attempt_count'] == 1) & ~(df['skill_name'].isnull())]

In [2]: response_df = pd.read_csv('correct.tsv', sep='\t').drop('Unnamed: 0', axis=1)
skill_df = pd.read_csv('skill.tsv', sep='\t').drop('Unnamed: 0', axis=1)
assistentment_df = pd.read_csv('assistentment_id.tsv', sep='\t').drop('Unnamed: 0', axis=1)
skill_dict = {}
with open('skill_dict.json', 'r', encoding='utf-8') as f:
    loaded = json.load(f)
    for k, v in loaded.items():
        skill_dict[k] = int(v)

skill_num = len(skill_dict) + 1 # including 0

def one_hot(skill_matrix, vocab_size):
    """
    params:
        skill_matrix: 2-D matrix (student, skills)
        vocab_size: size of the vocabulary
    returns:
        a ndarray with a shape like (student, sequence_len, vocab_size)
    """
    seq_len = skill_matrix.shape[1]
    result = np.zeros((skill_matrix.shape[0], seq_len, vocab_size))
    for i in range(skill_matrix.shape[0]):
        result[i, np.arange(seq_len), skill_matrix[i]] = 1.
    return result

def dkt_one_hot(skill_matrix, response_matrix, vocab_size):
```

```

seq_len = skill_matrix.shape[1]
skill_response_array = np.zeros((skill_matrix.shape[0], seq_len, 2 * vocab_size))
for i in range(skill_matrix.shape[0]):
    skill_response_array[i, np.arange(seq_len), 2 * skill_matrix[i] + response_matrix[i]] = response_matrix[i]
return skill_response_array

def preprocess(skill_df, response_df, skill_num):
    skill_matrix = skill_df.iloc[:, 1:].values
    response_array = response_df.iloc[:, 1:].values
    skill_array = one_hot(skill_matrix, skill_num)
    skill_response_array = dkt_one_hot(skill_matrix, response_array, skill_num)
    return skill_array, response_array, skill_response_array

```

```
skill_array, response_array, skill_response_array = preprocess(skill_df, response_df, skill_num)
```

```
In [38]: len(skill_dict)
```

```
Out[38]: 110
```

```
In [37]: response_array.shape
```

```
Out[37]: (584, 100)
```

```
In [26]: skill_array.shape
```

```
Out[26]: (584, 100, 111)
```

```
In [32]: skill_response_array.shape
```

```
Out[32]: (584, 100, 222)
```

```
In [3]: import keras
```

```

from keras.layers import Input, Dense, LSTM, TimeDistributed, Lambda, multiply
from keras.models import Model
from keras.optimizers import RMSprop, Adam
from keras.preprocessing.sequence import pad_sequences
from keras import backend as K

```

```

def build_skill2skill_model(input_shape, lstm_dim=32, dropout=0.0):
    input = Input(shape=input_shape, name='input skills')
    lstm = LSTM(lstm_dim,
                return_sequences=True,
                dropout=dropout,
                name='lstm layer')(input)
    output = TimeDistributed(Dense(input_shape[-1], activation='softmax'), name='probabilities')(lstm)
    model = Model(inputs=[input], outputs=[output])
    adam = Adam(lr=0.001, beta_1=0.9, beta_2=0.999, epsilon=1e-07, decay=0.0)
    model.compile(optimizer=adam, loss='categorical_crossentropy', metrics=['accuracy'])

```

```

model.summary()
return model

def reduce_dim(x):
    x = K.max(x, axis=-1, keepdims=True)
    return x

def build_dkt_model(input_shape, lstm_dim=32, dropout=0.0):
    input_skills = Input(shape=input_shape, name='input skills')
    lstm = LSTM(lstm_dim,
                return_sequences=True,
                dropout=dropout,
                name='lstm layer')(input_skills)
    dense = TimeDistributed(Dense(int(input_shape[-1]/2), activation='sigmoid'), name='p

    skill_next = Input(shape=(input_shape[0], int(input_shape[1]/2)), name='next_skill_t
    merged = multiply([dense, skill_next], name='multiply')
    reduced = Lambda(reduce_dim, output_shape=(input_shape[0], 1), name='reduce dim')(me

    model = Model(inputs=[input_skills, skill_next], outputs=[reduced])
    adam = Adam(lr=0.001, beta_1=0.9, beta_2=0.999, epsilon=1e-07, decay=0.0)
    model.compile(optimizer=adam, loss='binary_crossentropy', metrics=['accuracy'])
    model.summary()
    return model

print('skill2skill')
skill2skill_model = build_skill2skill_model((99, skill_num), lstm_dim=64)

print('dkt')
dkt_model = build_dkt_model((99, 2 * skill_num), lstm_dim=64)

```

Using Theano backend.

skill2skill

Layer (type)	Output Shape	Param #
input skills (InputLayer)	(None, 99, 111)	0
lstm layer (LSTM)	(None, 99, 64)	45056
probability (TimeDistributed)	(None, 99, 111)	7215

=====  
 Total params: 52,271  
 Trainable params: 52,271  
 Non-trainable params: 0

-----  
dkt

Layer (type)	Output Shape	Param #	Connected to
input skills (InputLayer)	(None, 99, 222)	0	
lstm layer (LSTM)	(None, 99, 64)	73472	input skills[0][0]
probability for each (TimeDistri	(None, 99, 111)	7215	lstm layer[0][0]
next_skill_tested (InputLayer)	(None, 99, 111)	0	
multiply (Multiply)	(None, 99, 111)	0	probability for each[0][0] next_skill_tested[0][0]
reduce dim (Lambda)	(None, 99, 1)	0	multiply[0][0]

=====

Total params: 80,687

Trainable params: 80,687

Non-trainable params: 0

-----

In [41]: %%time

```
# train skill2skill
skill2skill_model.fit(skill_array[:, 0:-1],
                      skill_array[:, 1:],
                      epochs=20,
                      batch_size=32,
                      shuffle=True,
                      validation_split=0.2)
```

Train on 467 samples, validate on 117 samples

Epoch 1/20

467/467 [=====] - 1s - loss: 4.6170 - acc: 0.1950 - val\_loss: 4.5729 -

Epoch 2/20

467/467 [=====] - 1s - loss: 4.2278 - acc: 0.1550 - val\_loss: 4.4697 -

Epoch 3/20

467/467 [=====] - 2s - loss: 3.6677 - acc: 0.1333 - val\_loss: 4.1492 -

Epoch 4/20

467/467 [=====] - 1s - loss: 3.1779 - acc: 0.2505 - val\_loss: 3.8165 -

Epoch 5/20

467/467 [=====] - 1s - loss: 2.7544 - acc: 0.3735 - val\_loss: 3.5075 -

Epoch 6/20

467/467 [=====] - 1s - loss: 2.4293 - acc: 0.4355 - val\_loss: 3.1725 -

Epoch 7/20

```

467/467 [=====] - 1s - loss: 2.1904 - acc: 0.4799 - val_loss: 2.9253 -
Epoch 8/20
467/467 [=====] - 1s - loss: 1.9651 - acc: 0.5337 - val_loss: 2.6795 -
Epoch 9/20
467/467 [=====] - 1s - loss: 1.7745 - acc: 0.5865 - val_loss: 2.4644 -
Epoch 10/20
467/467 [=====] - 1s - loss: 1.6159 - acc: 0.6326 - val_loss: 2.2535 -
Epoch 11/20
467/467 [=====] - 1s - loss: 1.5129 - acc: 0.6496 - val_loss: 2.1068 -
Epoch 12/20
467/467 [=====] - 1s - loss: 1.4089 - acc: 0.6717 - val_loss: 1.9672 -
Epoch 13/20
467/467 [=====] - 1s - loss: 1.2970 - acc: 0.6983 - val_loss: 1.8139 -
Epoch 14/20
467/467 [=====] - 1s - loss: 1.2062 - acc: 0.7132 - val_loss: 1.7264 -
Epoch 15/20
467/467 [=====] - 1s - loss: 1.1404 - acc: 0.7263 - val_loss: 1.6623 -
Epoch 16/20
467/467 [=====] - 1s - loss: 1.0881 - acc: 0.7363 - val_loss: 1.5911 -
Epoch 17/20
467/467 [=====] - 1s - loss: 1.0495 - acc: 0.7428 - val_loss: 1.5140 -
Epoch 18/20
467/467 [=====] - 1s - loss: 0.9877 - acc: 0.7582 - val_loss: 1.4528 -
Epoch 19/20
467/467 [=====] - 1s - loss: 0.9410 - acc: 0.7684 - val_loss: 1.4058 -
Epoch 20/20
467/467 [=====] - 1s - loss: 0.9027 - acc: 0.7805 - val_loss: 1.3694 -
CPU times: user 1min 15s, sys: 5.39 s, total: 1min 20s
Wall time: 2min 24s

```

Out[41]: <keras.callbacks.History at 0x129ea1d68>

In [42]: %%time

```

dkt_model.fit([skill_response_array[:, 0:-1], skill_array[:, 1:]],
              response_array[:, 1:, np.newaxis],
              epochs=20,
              batch_size=32,
              shuffle=True,
              validation_split=0.2)

```

Train on 467 samples, validate on 117 samples

```

Epoch 1/20
467/467 [=====] - 2s - loss: 0.6735 - acc: 0.7242 - val_loss: 0.6643 -
Epoch 2/20
467/467 [=====] - 3s - loss: 0.6072 - acc: 0.8360 - val_loss: 0.6066 -
Epoch 3/20

```

```

467/467 [=====] - 3s - loss: 0.4875 - acc: 0.8307 - val_loss: 0.5034 -
Epoch 4/20
467/467 [=====] - 2s - loss: 0.4167 - acc: 0.8489 - val_loss: 0.4606 -
Epoch 5/20
467/467 [=====] - 3s - loss: 0.3928 - acc: 0.8507 - val_loss: 0.4536 -
Epoch 6/20
467/467 [=====] - 4s - loss: 0.3828 - acc: 0.8552 - val_loss: 0.4447 -
Epoch 7/20
467/467 [=====] - 3s - loss: 0.3782 - acc: 0.8559 - val_loss: 0.4394 -
Epoch 8/20
467/467 [=====] - 3s - loss: 0.3772 - acc: 0.8555 - val_loss: 0.4479 -
Epoch 9/20
467/467 [=====] - 3s - loss: 0.3725 - acc: 0.8574 - val_loss: 0.4288 -
Epoch 10/20
467/467 [=====] - 2s - loss: 0.3686 - acc: 0.8571 - val_loss: 0.4392 -
Epoch 11/20
467/467 [=====] - 2s - loss: 0.3661 - acc: 0.8568 - val_loss: 0.4359 -
Epoch 12/20
467/467 [=====] - 2s - loss: 0.3628 - acc: 0.8571 - val_loss: 0.4344 -
Epoch 13/20
467/467 [=====] - 2s - loss: 0.3598 - acc: 0.8571 - val_loss: 0.4293 -
Epoch 14/20
467/467 [=====] - 2s - loss: 0.3571 - acc: 0.8587 - val_loss: 0.4287 -
Epoch 15/20
467/467 [=====] - 1s - loss: 0.3557 - acc: 0.8585 - val_loss: 0.4220 -
Epoch 16/20
467/467 [=====] - 2s - loss: 0.3541 - acc: 0.8600 - val_loss: 0.4316 -
Epoch 17/20
467/467 [=====] - 2s - loss: 0.3525 - acc: 0.8601 - val_loss: 0.4249 -
Epoch 18/20
467/467 [=====] - 2s - loss: 0.3494 - acc: 0.8615 - val_loss: 0.4194 -
Epoch 19/20
467/467 [=====] - 2s - loss: 0.3480 - acc: 0.8619 - val_loss: 0.4243 -
Epoch 20/20
467/467 [=====] - 2s - loss: 0.3488 - acc: 0.8622 - val_loss: 0.4199 -
CPU times: user 1min 22s, sys: 7.04 s, total: 1min 29s
Wall time: 1min 31s

```

```
Out[42]: <keras.callbacks.History at 0x12c76ae48>
```

## 1.1 Question 1

What were the 5 most common and 5 least common skills in this dataset? What percentage of responses are associated with the most common skill?

```
In [4]: skillname_df = pd.DataFrame(list(skill_dict.items()), columns=['Name', 'ID']).set_index(
        skillname_df.head()
```

```
Out[4]:
```

	Name
ID	
39	Estimation
90	Greatest Common Factor
88	Equation Solving More Than Two Steps
74	Translations
43	Multiplication Whole Numbers

```
In [5]: skill_counts = (skill_df
    .iloc[:, 1:]
    .unstack()
    .value_counts()
    .rename('count')
    .to_frame()
    .join(skillname_df)
    )
```

```
In [6]: print('Top 5 skills:')
        skill_counts.head()
```

Top 5 skills:

```
Out[6]:
```

	count	Name
7	4579	Table
30	4379	Conversion of Fraction Decimals Percents
8	3466	Venn Diagram
2	3404	Circle Graph
33	2833	Ordering Fractions

```
In [7]: print('Bottom 5 skills:')
        skill_counts.tail()
```

Bottom 5 skills:

```
Out[7]:
```

	count	Name
83	3	Volume Rectangular Prism
84	2	Volume Sphere
109	2	Solving Inequalities
82	2	Volume Cylinder
80	1	Surface Area Cylinder

```
In [8]: print('Proportion of responses for most common skill:')
        skill_counts.iloc[0, 0] / skill_counts['count'].sum()
```

Proportion of responses for most common skill:

```
Out[8]: 0.078407534246575344
```

## 1.2 Question 2

Train the sequence prediction model using a randomly selected 70% (training set) of students' data and predict on the remaining 30% (test set). What was the overall accuracy of skill prediction in the test set? What were the top 5 hardest and easiest to predict skills? Describe the metric you chose to represent hard/easy prediction.

```
In [9]: from sklearn.model_selection import train_test_split
```

```
        X_train, X_test, y_train, y_test = train_test_split(skill_array[:, 0:-1], skill_array[:,
        X_train.shape, X_test.shape
```

```
Out[9]: ((408, 99, 111), (176, 99, 111))
```

```
In [12]: %%time
```

```
        # train skill2skill
        skill2skill_model.fit(X_train,
                               y_train,
                               epochs=20,
                               batch_size=32,
                               shuffle=True,
                               validation_split=0.2)
```

Train on 326 samples, validate on 82 samples

Epoch 1/20

326/326 [=====] - 1s - loss: 4.6533 - acc: 0.1423 - val\_loss: 4.5828 -

Epoch 2/20

326/326 [=====] - 1s - loss: 4.5181 - acc: 0.4430 - val\_loss: 4.3526 -

Epoch 3/20

326/326 [=====] - 1s - loss: 4.1315 - acc: 0.2106 - val\_loss: 3.8672 -

Epoch 4/20

326/326 [=====] - 1s - loss: 3.7409 - acc: 0.1613 - val\_loss: 3.5118 -

Epoch 5/20

326/326 [=====] - 1s - loss: 3.4378 - acc: 0.2185 - val\_loss: 3.2544 -

Epoch 6/20

326/326 [=====] - 1s - loss: 3.1585 - acc: 0.3196 - val\_loss: 2.9795 -

Epoch 7/20

326/326 [=====] - 1s - loss: 2.8959 - acc: 0.3604 - val\_loss: 2.7486 -

Epoch 8/20

326/326 [=====] - 1s - loss: 2.6670 - acc: 0.3920 - val\_loss: 2.5466 -

Epoch 9/20

326/326 [=====] - 1s - loss: 2.4612 - acc: 0.4405 - val\_loss: 2.3498 -

Epoch 10/20

326/326 [=====] - 1s - loss: 2.3162 - acc: 0.4738 - val\_loss: 2.2786 -

Epoch 11/20

326/326 [=====] - 1s - loss: 2.1886 - acc: 0.5107 - val\_loss: 2.0904 -

Epoch 12/20

326/326 [=====] - 1s - loss: 2.0268 - acc: 0.5426 - val\_loss: 1.9686 -



```

Epoch 13/20
326/326 [=====] - 1s - loss: 1.8983 - acc: 0.5743 - val_loss: 1.8420 -
Epoch 14/20
326/326 [=====] - 1s - loss: 1.7876 - acc: 0.5972 - val_loss: 1.7449 -
Epoch 15/20
326/326 [=====] - 1s - loss: 1.6878 - acc: 0.6126 - val_loss: 1.6440 -
Epoch 16/20
326/326 [=====] - 1s - loss: 1.5961 - acc: 0.6331 - val_loss: 1.5593 -
Epoch 17/20
326/326 [=====] - 1s - loss: 1.5161 - acc: 0.6513 - val_loss: 1.4921 -
Epoch 18/20
326/326 [=====] - 1s - loss: 1.4482 - acc: 0.6656 - val_loss: 1.4310 -
Epoch 19/20
326/326 [=====] - 1s - loss: 1.3890 - acc: 0.6787 - val_loss: 1.3727 -
Epoch 20/20
326/326 [=====] - 1s - loss: 1.3295 - acc: 0.6963 - val_loss: 1.3053 -
CPU times: user 55.4 s, sys: 3.58 s, total: 58.9 s
Wall time: 38 s

```

```
Out[12]: <keras.callbacks.History at 0x1322a4240>
```

```
In [20]: test_predictions = skill2skill_model.predict(X_test)
         test_predictions.shape
```

```
Out[20]: (176, 99, 111)
```

```
In [74]: def s2s_acc(true, predictions):
         assert true.shape == predictions.shape
         return (np.count_nonzero(true.argmax(axis=2) == predictions.argmax(axis=2))
                 / (true.shape[0] * true.shape[1]))

         print("Overall accuracy:")
         s2s_acc(y_test, test_predictions)
```

```
Overall accuracy:
```

```
Out[74]: 0.7007575757575758
```

```
In [67]: from sklearn.metrics import classification

         scores = classification.recall_score(y_test.argmax(axis=2).flatten(),
                                             test_predictions.argmax(axis=2).flatten(),
                                             average=None)
         score_df = pd.DataFrame({'scores': scores}, index=np.unique(y_test.argmax(axis=2)))

         print('Easiest skills to predict (highest recall):')
         (score_df
```

```

        .sort_values('scores', ascending=False)
        .join(skillname_df)
        .head()
    )

```

Easiest skills to predict (highest recall):

```

Out[67]:
      scores                                     Name
1    0.984906                               Box and Whisker
46   0.957399                               Square Root
68   0.922145    Addition and Subtraction Integers
7    0.915704                               Table
8    0.910093                               Venn Diagram

```

```

In [66]: print('Hardest skills to predict (lowest recall):')
print('Note that there were a lot of 0s so these are just 5 of the 0s')
(score_df
 .sort_values('scores', ascending=False)
 .join(skillname_df)
 .tail()
 )

```

Hardest skills to predict (lowest recall):

Note that there were a lot of 0s so these are just 5 of the 0s

```

Out[66]:
      scores                                     Name
20      0.0    Angles on Parallel Lines Cut by a Transversal
56      0.0                               Pattern Finding
59      0.0                               Algebraic Solving
60      0.0    Choose an Equation from Given Information
110     0.0    Solving Systems of Linear Equations by Graphing

```

### 1.3 Question 3

Modify parameters of the network to increase accuracy (e.g. number of hidden nodes, optimizer, number of RNN layers, number of epochs, creating a validation set and stopping training when the validation set accuracy decreases). What were your accuracy results with respect to the hyper parameters you tuned?

```

In [88]: print('skill2skill with 128 hidden nodes')
s2s_128_model = build_skill2skill_model((99, skill_num), lstm_dim=128)

```

skill2skill with dropout of 0.1

Layer (type)	Output Shape	Param #
input skills (InputLayer)	(None, 99, 111)	0

```

-----
lstm layer (LSTM)                (None, 99, 128)                122880
-----
probability (TimeDistributed (None, 99, 111)                14319
=====
Total params: 137,199
Trainable params: 137,199
Non-trainable params: 0
-----

```

```
In [89]: %%time
```

```

# train skill2skill
s2s_128_model.fit(X_train,
                  y_train,
                  epochs=20,
                  batch_size=32,
                  shuffle=True,
                  validation_split=0.2)

```

Train on 326 samples, validate on 82 samples

```

Epoch 1/20
326/326 [=====] - 2s - loss: 4.6279 - acc: 0.2381 - val_loss: 4.4789 -
Epoch 2/20
326/326 [=====] - 2s - loss: 4.2272 - acc: 0.3127 - val_loss: 3.7845 -
Epoch 3/20
326/326 [=====] - 2s - loss: 3.6545 - acc: 0.2389 - val_loss: 3.2747 -
Epoch 4/20
326/326 [=====] - 2s - loss: 3.1287 - acc: 0.2737 - val_loss: 2.8435 -
Epoch 5/20
326/326 [=====] - 2s - loss: 2.6708 - acc: 0.3595 - val_loss: 2.4688 -
Epoch 6/20
326/326 [=====] - 2s - loss: 2.3373 - acc: 0.4474 - val_loss: 2.2300 -
Epoch 7/20
326/326 [=====] - 2s - loss: 2.1123 - acc: 0.4891 - val_loss: 1.9797 -
Epoch 8/20
326/326 [=====] - 2s - loss: 1.8995 - acc: 0.5505 - val_loss: 1.8180 -
Epoch 9/20
326/326 [=====] - 4s - loss: 1.7225 - acc: 0.5860 - val_loss: 1.7010 -
Epoch 10/20
326/326 [=====] - 2s - loss: 1.5751 - acc: 0.6170 - val_loss: 1.5052 -
Epoch 11/20
326/326 [=====] - 2s - loss: 1.4406 - acc: 0.6469 - val_loss: 1.4081 -
Epoch 12/20
326/326 [=====] - 2s - loss: 1.3413 - acc: 0.6756 - val_loss: 1.3315 -
Epoch 13/20
326/326 [=====] - 2s - loss: 1.2550 - acc: 0.7009 - val_loss: 1.2336 -

```

```

Epoch 14/20
326/326 [=====] - 2s - loss: 1.1747 - acc: 0.7134 - val_loss: 1.1567 -
Epoch 15/20
326/326 [=====] - 2s - loss: 1.1067 - acc: 0.7376 - val_loss: 1.0925 -
Epoch 16/20
326/326 [=====] - 2s - loss: 1.0463 - acc: 0.7471 - val_loss: 1.0376 -
Epoch 17/20
326/326 [=====] - 2s - loss: 1.0091 - acc: 0.7534 - val_loss: 0.9971 -
Epoch 18/20
326/326 [=====] - 2s - loss: 0.9565 - acc: 0.7662 - val_loss: 0.9530 -
Epoch 19/20
326/326 [=====] - 2s - loss: 0.9159 - acc: 0.7746 - val_loss: 0.9171 -
Epoch 20/20
326/326 [=====] - 2s - loss: 0.8811 - acc: 0.7826 - val_loss: 0.8956 -
CPU times: user 1min 31s, sys: 6.09 s, total: 1min 37s
Wall time: 59.8 s

```

```
Out[89]: <keras.callbacks.History at 0x140e574a8>
```

```

In [90]: print('Attempts to set dropout=0.1 or lower number of hidden nodes decreased accuracy.')

        print('Accuracy with 128 hidden nodes:')
        s2s_acc(y_test, s2s_128_model.predict(X_test))

```

```

Attempts to set dropout=0.1 or lower number of hidden nodes decreased accuracy.
Accuracy with 128 hidden nodes:

```

```
Out[90]: 0.7765725436179982
```

## 1.4 Question 4

Train a performance prediction model (DKT) using the same 70/30% split and report the accuracy and AUC of prediction on the 30%

```

In [100]: X1_train, X1_test, X2_train, X2_test, y_train, y_test = train_test_split(
            skill_response_array[:, 0:-1],
            skill_array[:, 1:],
            response_array[:, 1:, np.newaxis],
            test_size=0.3
        )
        X1_train.shape, X1_test.shape

```

```
Out[100]: ((408, 99, 222), (176, 99, 222))
```

```

In [101]: %%time

            dkt_model.fit([X1_train, X2_train],

```

```
y_train,
epochs=15,
batch_size=32,
shuffle=True,
validation_split=0.2)
```

Train on 326 samples, validate on 82 samples

Epoch 1/20

326/326 [=====] - 1s - loss: 0.6832 - acc: 0.6557 - val\_loss: 0.6666 -

Epoch 2/20

326/326 [=====] - 1s - loss: 0.6502 - acc: 0.8279 - val\_loss: 0.6302 -

Epoch 3/20

326/326 [=====] - 1s - loss: 0.5897 - acc: 0.8363 - val\_loss: 0.5278 -

Epoch 4/20

326/326 [=====] - 1s - loss: 0.4948 - acc: 0.8177 - val\_loss: 0.4607 -

Epoch 5/20

326/326 [=====] - 1s - loss: 0.4367 - acc: 0.8409 - val\_loss: 0.4309 -

Epoch 6/20

326/326 [=====] - 1s - loss: 0.4118 - acc: 0.8418 - val\_loss: 0.4242 -

Epoch 7/20

326/326 [=====] - 1s - loss: 0.4024 - acc: 0.8450 - val\_loss: 0.4182 -

Epoch 8/20

326/326 [=====] - 1s - loss: 0.3949 - acc: 0.8475 - val\_loss: 0.4171 -

Epoch 9/20

326/326 [=====] - 1s - loss: 0.3932 - acc: 0.8507 - val\_loss: 0.4122 -

Epoch 10/20

326/326 [=====] - 1s - loss: 0.3854 - acc: 0.8511 - val\_loss: 0.4100 -

Epoch 11/20

326/326 [=====] - 1s - loss: 0.3806 - acc: 0.8513 - val\_loss: 0.4054 -

Epoch 12/20

326/326 [=====] - 1s - loss: 0.3781 - acc: 0.8528 - val\_loss: 0.4016 -

Epoch 13/20

326/326 [=====] - 1s - loss: 0.3753 - acc: 0.8529 - val\_loss: 0.3979 -

Epoch 14/20

326/326 [=====] - 1s - loss: 0.3709 - acc: 0.8545 - val\_loss: 0.3997 -

Epoch 15/20

326/326 [=====] - 1s - loss: 0.3667 - acc: 0.8557 - val\_loss: 0.3940 -

Epoch 16/20

326/326 [=====] - 1s - loss: 0.3635 - acc: 0.8568 - val\_loss: 0.3928 -

Epoch 17/20

326/326 [=====] - 1s - loss: 0.3597 - acc: 0.8587 - val\_loss: 0.3916 -

Epoch 18/20

326/326 [=====] - 1s - loss: 0.3583 - acc: 0.8590 - val\_loss: 0.3929 -

Epoch 19/20

326/326 [=====] - 1s - loss: 0.3563 - acc: 0.8590 - val\_loss: 0.3898 -

Epoch 20/20

326/326 [=====] - 1s - loss: 0.3568 - acc: 0.8599 - val\_loss: 0.3877 -

CPU times: user 52.9 s, sys: 3.69 s, total: 56.6 s

Wall time: 34.9 s

Out[101]: <keras.callbacks.History at 0x16af02390>

In [119]: `dkt_model.evaluate([X1_test, X2_test], y_test)`

176/176 [=====] - 0s

Out[119]: [0.35878889127211139, 0.86346419291062793]

In [121]: `from sklearn.metrics import accuracy_score`

```
dkt_predictions = dk_model.predict([X1_test, X2_test])
print('DKT Prediction Accuracy:')
accuracy_score(np.round(dkt_predictions.flatten()), y_test.flatten())
```

DKT Prediction Accuracy:

Out[121]: 0.86346418732782371

In [140]: `from sklearn.metrics import roc_auc_score`

```
print('DKT AUC score:')
roc_auc_score(y_test.flatten(), dk_model.predict([X1_test, X2_test]).flatten())
```

DKT AUC score:

Out[140]: 0.74280793994244021

## 1.5 Question 5

Tune the hyper parameters of this model to improve accuracy and report your improvement with respect to the tuned parameters. Which lead to the most significant improvement?

```
In [124]: print('dkt with 128 hidden nodes')
          dk_model_128 = build_dkt_model((99, 2 * skill_num), lstm_dim=128)
```

dkt with 128 hidden nodes

Layer (type)	Output Shape	Param #	Connected to
input skills (InputLayer)	(None, 99, 222)	0	
lstm layer (LSTM)	(None, 99, 128)	179712	input skills[0][0]
probability for each (TimeDistri	(None, 99, 111)	14319	lstm layer[0][0]

```

-----
next_skill_tested (InputLayer)    (None, 99, 111)    0
-----
multiply (Multiply)                (None, 99, 111)    0    probability for each[0][0]
                                   next_skill_tested[0][0]
-----
reduce dim (Lambda)                (None, 99, 1)      0    multiply[0][0]
=====
Total params: 194,031
Trainable params: 194,031
Non-trainable params: 0
-----

```

In [125]: %%time

```

dkt_128_model.fit([X1_train, X2_train],
                  y_train,
                  epochs=20,
                  batch_size=32,
                  shuffle=True,
                  validation_split=0.2)

```

Train on 326 samples, validate on 82 samples

```

Epoch 1/20
326/326 [=====] - 2s - loss: 0.6695 - acc: 0.7454 - val_loss: 0.6366 -
Epoch 2/20
326/326 [=====] - 2s - loss: 0.5696 - acc: 0.8423 - val_loss: 0.4805 -
Epoch 3/20
326/326 [=====] - 2s - loss: 0.4444 - acc: 0.8397 - val_loss: 0.4346 -
Epoch 4/20
326/326 [=====] - 2s - loss: 0.4114 - acc: 0.8429 - val_loss: 0.4258 -
Epoch 5/20
326/326 [=====] - 2s - loss: 0.4038 - acc: 0.8452 - val_loss: 0.4159 -
Epoch 6/20
326/326 [=====] - 2s - loss: 0.3959 - acc: 0.8491 - val_loss: 0.4138 -
Epoch 7/20
326/326 [=====] - 2s - loss: 0.3937 - acc: 0.8488 - val_loss: 0.4104 -
Epoch 8/20
326/326 [=====] - 2s - loss: 0.3882 - acc: 0.8513 - val_loss: 0.4072 -
Epoch 9/20
326/326 [=====] - 2s - loss: 0.3842 - acc: 0.8500 - val_loss: 0.4052 -
Epoch 10/20
326/326 [=====] - 2s - loss: 0.3948 - acc: 0.8472 - val_loss: 0.4102 -
Epoch 11/20
326/326 [=====] - 2s - loss: 0.3878 - acc: 0.8480 - val_loss: 0.4080 -
Epoch 12/20
326/326 [=====] - 2s - loss: 0.3861 - acc: 0.8524 - val_loss: 0.4085 -

```

```

Epoch 13/20
326/326 [=====] - 2s - loss: 0.3791 - acc: 0.8534 - val_loss: 0.4025 -
Epoch 14/20
326/326 [=====] - 2s - loss: 0.3763 - acc: 0.8520 - val_loss: 0.3985 -
Epoch 15/20
326/326 [=====] - 2s - loss: 0.3708 - acc: 0.8558 - val_loss: 0.3971 -
Epoch 16/20
326/326 [=====] - 2s - loss: 0.3703 - acc: 0.8545 - val_loss: 0.3967 -
Epoch 17/20
326/326 [=====] - 3s - loss: 0.3650 - acc: 0.8575 - val_loss: 0.3975 -
Epoch 18/20
326/326 [=====] - 2s - loss: 0.3670 - acc: 0.8578 - val_loss: 0.3931 -
Epoch 19/20
326/326 [=====] - 2s - loss: 0.3612 - acc: 0.8576 - val_loss: 0.3919 -
Epoch 20/20
326/326 [=====] - 2s - loss: 0.3598 - acc: 0.8570 - val_loss: 0.3868 -
CPU times: user 1min 26s, sys: 5.34 s, total: 1min 31s
Wall time: 53.6 s

```

```
Out[125]: <keras.callbacks.History at 0x144d1b7f0>
```

```

In [126]: dkt_128_predictions = dkt_128_model.predict([X1_test, X2_test])
          print('DKT Prediction Accuracy with 128 hidden nodes:')
          accuracy_score(np.round(dkt_128_predictions.flatten()), y_test.flatten())

```

DKT Prediction Accuracy with 128 hidden nodes:

```
Out[126]: 0.86116850321395777
```

That didn't really work; attempt 2:

```

In [128]: print('dkt with 128 hidden nodes and dropout of 0.01')
          dkt_128_model = build_dkt_model((99, 2 * skill_num), lstm_dim=128, dropout=0.01)

```

dkt with 128 hidden nodes and dropout of 0.01

Layer (type)	Output Shape	Param #	Connected to
input skills (InputLayer)	(None, 99, 222)	0	
lstm layer (LSTM)	(None, 99, 128)	179712	input skills[0][0]
probability for each (TimeDistri	(None, 99, 111)	14319	lstm layer[0][0]
next_skill_tested (InputLayer)	(None, 99, 111)	0	
multiply (Multiply)	(None, 99, 111)	0	probability for each[0][0]



```

next_skill_tested[0][0]
-----
reduce dim (Lambda)          (None, 99, 1)          0          multiply[0][0]
=====
Total params: 194,031
Trainable params: 194,031
Non-trainable params: 0
-----

```

```
In [129]: %%time
```

```

dkt_128_model.fit([X1_train, X2_train],
                  y_train,
                  epochs=20,
                  batch_size=32,
                  shuffle=True,
                  validation_split=0.2)

```

Train on 326 samples, validate on 82 samples

```

Epoch 1/20
326/326 [=====] - 2s - loss: 0.6687 - acc: 0.7452 - val_loss: 0.6373 -
Epoch 2/20
326/326 [=====] - 2s - loss: 0.5703 - acc: 0.8359 - val_loss: 0.5011 -
Epoch 3/20
326/326 [=====] - 2s - loss: 0.4598 - acc: 0.8313 - val_loss: 0.4398 -
Epoch 4/20
326/326 [=====] - 2s - loss: 0.4196 - acc: 0.8409 - val_loss: 0.4253 -
Epoch 5/20
326/326 [=====] - 2s - loss: 0.4059 - acc: 0.8445 - val_loss: 0.4191 -
Epoch 6/20
326/326 [=====] - 2s - loss: 0.3997 - acc: 0.8461 - val_loss: 0.4172 -
Epoch 7/20
326/326 [=====] - 2s - loss: 0.3960 - acc: 0.8486 - val_loss: 0.4121 -
Epoch 8/20
326/326 [=====] - 2s - loss: 0.3903 - acc: 0.8488 - val_loss: 0.4089 -
Epoch 9/20
326/326 [=====] - 2s - loss: 0.3865 - acc: 0.8498 - val_loss: 0.4077 -
Epoch 10/20
326/326 [=====] - 2s - loss: 0.3842 - acc: 0.8519 - val_loss: 0.4054 -
Epoch 11/20
326/326 [=====] - 2s - loss: 0.3800 - acc: 0.8533 - val_loss: 0.4020 -
Epoch 12/20
326/326 [=====] - 2s - loss: 0.3765 - acc: 0.8519 - val_loss: 0.4003 -
Epoch 13/20
326/326 [=====] - 2s - loss: 0.3731 - acc: 0.8549 - val_loss: 0.3970 -
Epoch 14/20
326/326 [=====] - 2s - loss: 0.3724 - acc: 0.8534 - val_loss: 0.4017 -

```

```

Epoch 15/20
326/326 [=====] - 2s - loss: 0.3773 - acc: 0.8450 - val_loss: 0.3988 -
Epoch 16/20
326/326 [=====] - 3s - loss: 0.3728 - acc: 0.8470 - val_loss: 0.3920 -
Epoch 17/20
326/326 [=====] - 3s - loss: 0.3681 - acc: 0.8491 - val_loss: 0.3939 -
Epoch 18/20
326/326 [=====] - 2s - loss: 0.3628 - acc: 0.8535 - val_loss: 0.3935 -
Epoch 19/20
326/326 [=====] - 2s - loss: 0.3593 - acc: 0.8573 - val_loss: 0.3844 -
Epoch 20/20
326/326 [=====] - 2s - loss: 0.3530 - acc: 0.8610 - val_loss: 0.3838 -
CPU times: user 1min 34s, sys: 8.24 s, total: 1min 43s
Wall time: 60 s

```

```
Out[129]: <keras.callbacks.History at 0x147370278>
```

```

In [130]: dkt_128_predictions = dkt_128_model.predict([X1_test, X2_test])
          print('DKT Prediction Accuracy with 128 hidden nodes and dropout of 0.01:')
          accuracy_score(np.round(dkt_128_predictions.flatten()), y_test.flatten())

```

DKT Prediction Accuracy with 128 hidden nodes and dropout of 0.01:

```
Out[130]: 0.8617424242424242
```

```

In [139]: print('DKT AUC score:')
          roc_auc_score(y_test.flatten(), dkt_128_predictions.flatten())

```

DKT AUC score:

```
Out[139]: 0.75305468559654443
```

That was about the same as well :(.

```

In [132]: print('dkt with 256 hidden nodes and dropout of 0.01')
          dkt_256_model = build_dkt_model((99, 2 * skill_num), lstm_dim=256, dropout=0.01)

```

dkt with 256 hidden nodes and dropout of 0.01

Layer (type)	Output Shape	Param #	Connected to
input skills (InputLayer)	(None, 99, 222)	0	
lstm layer (LSTM)	(None, 99, 256)	490496	input skills[0][0]
probability for each (TimeDistri	(None, 99, 111)	28527	lstm layer[0][0]

```

-----
next_skill_tested (InputLayer)  (None, 99, 111)      0
-----
multiply (Multiply)             (None, 99, 111)      0      probability for each[0][0]
                                   next_skill_tested[0][0]
-----
reduce dim (Lambda)            (None, 99, 1)        0      multiply[0][0]
=====
Total params: 519,023
Trainable params: 519,023
Non-trainable params: 0
-----

```

```
In [133]: %%time
```

```

    dkt_256_model.fit([X1_train, X2_train],
                      y_train,
                      epochs=20,
                      batch_size=32,
                      shuffle=True,
                      validation_split=0.2)

```

Train on 326 samples, validate on 82 samples

```

Epoch 1/20
326/326 [=====] - 5s - loss: 0.6452 - acc: 0.7807 - val_loss: 0.5479 -
Epoch 2/20
326/326 [=====] - 5s - loss: 0.4646 - acc: 0.8382 - val_loss: 0.4357 -
Epoch 3/20
326/326 [=====] - 5s - loss: 0.4175 - acc: 0.8417 - val_loss: 0.4228 -
Epoch 4/20
326/326 [=====] - 5s - loss: 0.4064 - acc: 0.8443 - val_loss: 0.4128 -
Epoch 5/20
326/326 [=====] - 5s - loss: 0.3960 - acc: 0.8454 - val_loss: 0.4106 -
Epoch 6/20
326/326 [=====] - 5s - loss: 0.3923 - acc: 0.8478 - val_loss: 0.4106 -
Epoch 7/20
326/326 [=====] - 5s - loss: 0.3878 - acc: 0.8524 - val_loss: 0.4125 -
Epoch 8/20
326/326 [=====] - 5s - loss: 0.3875 - acc: 0.8468 - val_loss: 0.4039 -
Epoch 9/20
326/326 [=====] - 5s - loss: 0.3854 - acc: 0.8490 - val_loss: 0.3999 -
Epoch 10/20
326/326 [=====] - 5s - loss: 0.3779 - acc: 0.8543 - val_loss: 0.4008 -
Epoch 11/20
326/326 [=====] - 5s - loss: 0.3749 - acc: 0.8569 - val_loss: 0.3961 -
Epoch 12/20
326/326 [=====] - 5s - loss: 0.3679 - acc: 0.8591 - val_loss: 0.3938 -

```

```

Epoch 13/20
326/326 [=====] - 5s - loss: 0.3628 - acc: 0.8602 - val_loss: 0.3900 -
Epoch 14/20
326/326 [=====] - 5s - loss: 0.3584 - acc: 0.8623 - val_loss: 0.3988 -
Epoch 15/20
326/326 [=====] - 5s - loss: 0.3597 - acc: 0.8617 - val_loss: 0.3868 -
Epoch 16/20
326/326 [=====] - 5s - loss: 0.3570 - acc: 0.8621 - val_loss: 0.3835 -
Epoch 17/20
326/326 [=====] - 5s - loss: 0.3524 - acc: 0.8625 - val_loss: 0.3843 -
Epoch 18/20
326/326 [=====] - 5s - loss: 0.3519 - acc: 0.8625 - val_loss: 0.3849 -
Epoch 19/20
326/326 [=====] - 5s - loss: 0.3494 - acc: 0.8636 - val_loss: 0.3777 -
Epoch 20/20
326/326 [=====] - 5s - loss: 0.3447 - acc: 0.8643 - val_loss: 0.3779 -
CPU times: user 3min 24s, sys: 13 s, total: 3min 37s
Wall time: 1min 57s

```

```
Out[133]: <keras.callbacks.History at 0x141d05080>
```

```

In [134]: dkt_256_predictions = dkt_256_model.predict([X1_test, X2_test])
          print('DKT Prediction Accuracy with 256 hidden nodes and dropout of 0.01:')
          accuracy_score(np.round(dkt_256_predictions.flatten()), y_test.flatten())

```

DKT Prediction Accuracy with 256 hidden nodes and dropout of 0.01:

```
Out[134]: 0.8621441689623508
```

```

In [137]: print('DKT AUC score:')
          roc_auc_score(y_test.flatten(), dkt_256_predictions.flatten())

```

DKT AUC score:

```
Out[137]: 0.75923135423709032
```

**Looks like we were able to get the AUC to increase by 1.7% compared to the original model using 256 hidden nodes and a dropout of 0.1.**