

dk

September 26, 2017

1 RNN modeling of behavior and performance

Yifat Amir, Allen Guo, Sam Lau

```
In [1]: import pandas as pd
import numpy as np
import json

filename = 'skill_builder_data_corrected.csv'
df = pd.read_csv(filename, encoding='ISO-8859-1', low_memory=False)
df = df[(df['original'] == 1) & (df['attempt_count'] == 1) & ~(df['skill_name'].isnull())]

In [2]: response_df = pd.read_csv('correct.tsv', sep='\t').drop('Unnamed: 0', axis=1)
skill_df = pd.read_csv('skill.tsv', sep='\t').drop('Unnamed: 0', axis=1)
assistentment_df = pd.read_csv('assistentment_id.tsv', sep='\t').drop('Unnamed: 0', axis=1)
skill_dict = {}
with open('skill_dict.json', 'r', encoding='utf-8') as f:
    loaded = json.load(f)
    for k, v in loaded.items():
        skill_dict[k] = int(v)

skill_num = len(skill_dict) + 1 # including 0

def one_hot(skill_matrix, vocab_size):
    """
    params:
        skill_matrix: 2-D matrix (student, skills)
        vocab_size: size of the vocabulary
    returns:
        a ndarray with a shape like (student, sequence_len, vocab_size)
    """
    seq_len = skill_matrix.shape[1]
    result = np.zeros((skill_matrix.shape[0], seq_len, vocab_size))
    for i in range(skill_matrix.shape[0]):
        result[i, np.arange(seq_len), skill_matrix[i]] = 1.
    return result

def dkt_one_hot(skill_matrix, response_matrix, vocab_size):
```

```

seq_len = skill_matrix.shape[1]
skill_response_array = np.zeros((skill_matrix.shape[0], seq_len, 2 * vocab_size))
for i in range(skill_matrix.shape[0]):
    skill_response_array[i, np.arange(seq_len), 2 * skill_matrix[i] + response_matrix[i]] = response_matrix[i]
return skill_response_array

def preprocess(skill_df, response_df, skill_num):
    skill_matrix = skill_df.iloc[:, 1:].values
    response_array = response_df.iloc[:, 1:].values
    skill_array = one_hot(skill_matrix, skill_num)
    skill_response_array = dkt_one_hot(skill_matrix, response_array, skill_num)
    return skill_array, response_array, skill_response_array

```

```

skill_array, response_array, skill_response_array = preprocess(skill_df, response_df, skill_num)

```

```

In [ ]: import keras
from keras.layers import Input, Dense, LSTM, TimeDistributed, Lambda, multiply
from keras.models import Model
from keras.optimizers import RMSprop, Adam
from keras.preprocessing.sequence import pad_sequences
from keras import backend as K

def build_skill2skill_model(input_shape, lstm_dim=32, dropout=0.0):
    input = Input(shape=input_shape, name='input skills')
    lstm = LSTM(lstm_dim,
                return_sequences=True,
                dropout=dropout,
                name='lstm layer')(input)
    output = TimeDistributed(Dense(input_shape[-1], activation='softmax'), name='probabilities')(lstm)
    model = Model(inputs=[input], outputs=[output])
    adam = Adam(lr=0.001, beta_1=0.9, beta_2=0.999, epsilon=1e-07, decay=0.0)
    model.compile(optimizer=adam, loss='categorical_crossentropy', metrics=['accuracy'])
    model.summary()
    return model

def reduce_dim(x):
    x = K.max(x, axis=-1, keepdims=True)
    return x

def build_dkt_model(input_shape, lstm_dim=32, dropout=0.0):
    input_skills = Input(shape=input_shape, name='input skills')
    lstm = LSTM(lstm_dim,
                return_sequences=True,
                dropout=dropout,
                name='lstm layer')(input_skills)
    dense = TimeDistributed(Dense(int(input_shape[-1]/2), activation='sigmoid'), name='probabilities')(lstm)
    model = Model(inputs=[input_skills], outputs=[dense])
    adam = Adam(lr=0.001, beta_1=0.9, beta_2=0.999, epsilon=1e-07, decay=0.0)
    model.compile(optimizer=adam, loss='binary_crossentropy', metrics=['accuracy'])
    model.summary()
    return model

```

```

skill_next = Input(shape=(input_shape[0], int(input_shape[1]/2)), name='next_skill_t')
merged = multiply([dense, skill_next], name='multiply')
reduced = Lambda(reduce_dim, output_shape=(input_shape[0], 1), name='reduce_dim')(merged)

model = Model(inputs=[input_skills, skill_next], outputs=[reduced])
adam = Adam(lr=0.001, beta_1=0.9, beta_2=0.999, epsilon=1e-07, decay=0.0)
model.compile(optimizer=adam, loss='binary_crossentropy', metrics=['accuracy'])
model.summary()
return model

print('skill2skill')
skill2skill_model = build_skill2skill_model((99, skill_num), lstm_dim=64)

print('dkt')
dkt_model = build_dkt_model((99, 2 * skill_num), lstm_dim=64)

```

```
In [ ]: %%time
```

```

# train skill2skill
skill2skill_model.fit(skill_array[:, 0:-1],
                      skill_array[:, 1:],
                      epochs=20,
                      batch_size=32,
                      shuffle=True,
                      validation_split=0.2)

```

```
In [ ]: %%time
```

```

dkt_model.fit([skill_response_array[:, 0:-1], skill_array[:, 1:]],
              response_array[:, 1:, np.newaxis],
              epochs=20,
              batch_size=32,
              shuffle=True,
              validation_split=0.2)

```

1.1 Question 1

What were the 5 most common and 5 least common skills in this dataset? What percentage of responses are associated with the most common skill?

```
In [4]: skillname_df = pd.DataFrame(list(skill_dict.items()), columns=['Name', 'ID']).set_index('ID')
skillname_df.head()
```

```
Out[4]:
```

	Name
ID	
39	Estimation
90	Greatest Common Factor
88	Equation Solving More Than Two Steps

74 Translations
 43 Multiplication Whole Numbers

```
In [5]: skill_counts = (skill_df
    .iloc[:, 1:]
    .unstack()
    .value_counts()
    .rename('count')
    .to_frame()
    .join(skillname_df)
    )
```

```
In [6]: print('Top 5 skills:')
    skill_counts.head()
```

Top 5 skills:

```
Out[6]:
```

	count	Name
7	4579	Table
30	4379	Conversion of Fraction Decimals Percents
8	3466	Venn Diagram
2	3404	Circle Graph
33	2833	Ordering Fractions

```
In [7]: print('Bottom 5 skills:')
    skill_counts.tail()
```

Bottom 5 skills:

```
Out[7]:
```

	count	Name
83	3	Volume Rectangular Prism
84	2	Volume Sphere
109	2	Solving Inequalities
82	2	Volume Cylinder
80	1	Surface Area Cylinder

```
In [8]: print('Proportion of responses for most common skill:')
    skill_counts.iloc[0, 0] / skill_counts['count'].sum()
```

Proportion of responses for most common skill:

```
Out[8]: 0.078407534246575344
```

1.2 Question 2

Train the sequence prediction model using a randomly selected 70% (training set) of students' data and predict on the remaining 30% (test set). What was the overall accuracy of skill prediction in the test set? What were the top 5 hardest and easiest to predict skills? Describe the metric you chose to represent hard/easy prediction.

```
In [9]: from sklearn.model_selection import train_test_split
```

```
    X_train, X_test, y_train, y_test = train_test_split(skill_array[:, 0:-1], skill_array[:, 0],
    X_train.shape, X_test.shape)
```

```
Out[9]: ((408, 99, 111), (176, 99, 111))
```

```
In [ ]: %%time
```

```
    # train skill2skill
    skill2skill_model.fit(X_train,
                           y_train,
                           epochs=20,
                           batch_size=32,
                           shuffle=True,
                           validation_split=0.2)
```

```
In [20]: test_predictions = skill2skill_model.predict(X_test)
    test_predictions.shape
```

```
Out[20]: (176, 99, 111)
```

```
In [74]: def s2s_acc(true, predictions):
    assert true.shape == predictions.shape
    return (np.count_nonzero(true.argmax(axis=2) == predictions.argmax(axis=2))
            / (true.shape[0] * true.shape[1]))

    print("Overall accuracy:")
    s2s_acc(y_test, test_predictions)
```

Overall accuracy:

```
Out[74]: 0.7007575757575758
```

```
In [67]: from sklearn.metrics import classification
```

```
    scores = classification.recall_score(y_test.argmax(axis=2).flatten(),
                                         test_predictions.argmax(axis=2).flatten(),
                                         average=None)
    score_df = pd.DataFrame({'scores': scores}, index=np.unique(y_test.argmax(axis=2)))

    print('Easiest skills to predict (highest recall):')
    (score_df
     .sort_values('scores', ascending=False)
     .join(skillname_df)
     .head()
    )
```

Easiest skills to predict (highest recall):

```
Out[67]:
```

	scores	Name
1	0.984906	Box and Whisker
46	0.957399	Square Root
68	0.922145	Addition and Subtraction Integers
7	0.915704	Table
8	0.910093	Venn Diagram

```
In [66]: print('Hardest skills to predict (lowest recall):')
print('Note that there were a lot of 0s so these are just 5 of the 0s')
(score_df
 .sort_values('scores', ascending=False)
 .join(skillname_df)
 .tail()
 )
```

Hardest skills to predict (lowest recall):

Note that there were a lot of 0s so these are just 5 of the 0s

```
Out[66]:
```

	scores	Name
20	0.0	Angles on Parallel Lines Cut by a Transversal
56	0.0	Pattern Finding
59	0.0	Algebraic Solving
60	0.0	Choose an Equation from Given Information
110	0.0	Solving Systems of Linear Equations by Graphing

1.3 Question 3

Modify parameters of the network to increase accuracy (e.g. number of hidden nodes, optimizer, number of RNN layers, number of epochs, creating a validation set and stopping training when the validation set accuracy decreases). What were your accuracy results with respect to the hyper parameters you tuned?

```
In [88]: print('skill2skill with 128 hidden nodes')
s2s_128_model = build_skill2skill_model((99, skill_num), lstm_dim=128)
```

skill2skill with dropout of 0.1

Layer (type)	Output Shape	Param #
input skills (InputLayer)	(None, 99, 111)	0
lstm layer (LSTM)	(None, 99, 128)	122880
probability (TimeDistributed	(None, 99, 111)	14319

Total params: 137,199
Trainable params: 137,199
Non-trainable params: 0

```
In [ ]: %%time
```

```
# train skill2skill
s2s_128_model.fit(X_train,
                  y_train,
                  epochs=20,
                  batch_size=32,
                  shuffle=True,
                  validation_split=0.2)
```

```
In [90]: print('Attempts to set dropout=0.1 or lower number of hidden nodes decreased accuracy.'
```

```
print('Accuracy with 128 hidden nodes:')
s2s_acc(y_test, s2s_128_model.predict(X_test))
```

Attempts to set dropout=0.1 or lower number of hidden nodes decreased accuracy.
Accuracy with 128 hidden nodes:

```
Out[90]: 0.7765725436179982
```

1.4 Question 4

Train a performance prediction model (DKT) using the same 70/30% split and report the accuracy and AUC of prediction on the 30%

```
In [100]: X1_train, X1_test, X2_train, X2_test, y_train, y_test = train_test_split(
            skill_response_array[:, 0:-1],
            skill_array[:, 1:],
            response_array[:, 1:, np.newaxis],
            test_size=0.3
        )
        X1_train.shape, X1_test.shape
```

```
Out[100]: ((408, 99, 222), (176, 99, 222))
```

```
In [ ]: %%time
```

```
dkt_model.fit([X1_train, X2_train],
              y_train,
              epochs=15,
              batch_size=32,
              shuffle=True,
              validation_split=0.2)
```

```
In [119]: dkt_model.evaluate([X1_test, X2_test], y_test)
```

```
176/176 [=====] - 0s
```

```
Out[119]: [0.35878889127211139, 0.86346419291062793]
```

```
In [121]: from sklearn.metrics import accuracy_score
```

```
    dkt_predictions = dkt_model.predict([X1_test, X2_test])
    print('DKT Prediction Accuracy:')
    accuracy_score(np.round(dkt_predictions.flatten()), y_test.flatten())
```

DKT Prediction Accuracy:

```
Out[121]: 0.86346418732782371
```

```
In [140]: from sklearn.metrics import roc_auc_score
```

```
    print('DKT AUC score:')
    roc_auc_score(y_test.flatten(), dkt_predictions.flatten())
```

DKT AUC score:

```
Out[140]: 0.74280793994244021
```

1.5 Question 5

Tune the hyper parameters of this model to improve accuracy and report your improvement with respect to the tuned parameters. Which lead to the most significant improvement?

```
In [124]: print('dkt with 128 hidden nodes')
          dkt_128_model = build_dkt_model((99, 2 * skill_num), lstm_dim=128)
```

dkt with 128 hidden nodes

Layer (type)	Output Shape	Param #	Connected to
input skills (InputLayer)	(None, 99, 222)	0	
lstm layer (LSTM)	(None, 99, 128)	179712	input skills[0][0]
probability for each (TimeDistri	(None, 99, 111)	14319	lstm layer[0][0]
next_skill_tested (InputLayer)	(None, 99, 111)	0	
multiply (Multiply)	(None, 99, 111)	0	probability for each[0][0] next_skill_tested[0][0]

reduce dim (Lambda)	(None, 99, 1)	0	multiply[0][0]
---------------------	---------------	---	----------------

```

Total params: 194,031
Trainable params: 194,031
Non-trainable params: 0

```

```
In [ ]: %%time
```

```

dkt_128_model.fit([X1_train, X2_train],
                  y_train,
                  epochs=20,
                  batch_size=32,
                  shuffle=True,
                  validation_split=0.2)

```

```

In [126]: dkt_128_predictions = dkt_128_model.predict([X1_test, X2_test])
          print('DKT Prediction Accuracy with 128 hidden nodes:')
          accuracy_score(np.round(dkt_128_predictions.flatten()), y_test.flatten())

```

DKT Prediction Accuracy with 128 hidden nodes:

```
Out[126]: 0.86116850321395777
```

That didn't really work; attempt 2:

```

In [128]: print('dkt with 128 hidden nodes and dropout of 0.01')
          dkt_128_model = build_dkt_model((99, 2 * skill_num), lstm_dim=128, dropout=0.01)

```

dkt with 128 hidden nodes and dropout of 0.01

Layer (type)	Output Shape	Param #	Connected to
input skills (InputLayer)	(None, 99, 222)	0	
lstm layer (LSTM)	(None, 99, 128)	179712	input skills[0][0]
probability for each (TimeDistri	(None, 99, 111)	14319	lstm layer[0][0]
next_skill_tested (InputLayer)	(None, 99, 111)	0	
multiply (Multiply)	(None, 99, 111)	0	probability for each[0][0] next_skill_tested[0][0]
reduce dim (Lambda)	(None, 99, 1)	0	multiply[0][0]

Total params: 194,031
Trainable params: 194,031
Non-trainable params: 0

In []: %%time

```
dkt_128_model.fit([X1_train, X2_train],
                  y_train,
                  epochs=20,
                  batch_size=32,
                  shuffle=True,
                  validation_split=0.2)
```

```
In [130]: dk128_predictions = dk128_model.predict([X1_test, X2_test])
          print('DKT Prediction Accuracy with 128 hidden nodes and dropout of 0.01:')
          accuracy_score(np.round(dk128_predictions.flatten()), y_test.flatten())
```

DKT Prediction Accuracy with 128 hidden nodes and dropout of 0.01:

Out[130]: 0.8617424242424242

```
In [139]: print('DKT AUC score:')
          roc_auc_score(y_test.flatten(), dk128_predictions.flatten())
```

DKT AUC score:

Out[139]: 0.75305468559654443

That was about the same as well :(.

```
In [132]: print('dkt with 256 hidden nodes and dropout of 0.01')
          dk256_model = build_dkt_model((99, 2 * skill_num), lstm_dim=256, dropout=0.01)
```

dkt with 256 hidden nodes and dropout of 0.01

Layer (type)	Output Shape	Param #	Connected to
input skills (InputLayer)	(None, 99, 222)	0	
lstm layer (LSTM)	(None, 99, 256)	490496	input skills[0][0]
probability for each (TimeDistri	(None, 99, 111)	28527	lstm layer[0][0]
next_skill_tested (InputLayer)	(None, 99, 111)	0	

multiply (Multiply)	(None, 99, 111)	0	probability for each[0][0] next_skill_tested[0][0]
---------------------	-----------------	---	---

reduce dim (Lambda)	(None, 99, 1)	0	multiply[0][0]
---------------------	---------------	---	----------------

```

Total params: 519,023
Trainable params: 519,023
Non-trainable params: 0

```

```
In [ ]: %%time
```

```

dkt_256_model.fit([X1_train, X2_train],
                  y_train,
                  epochs=20,
                  batch_size=32,
                  shuffle=True,
                  validation_split=0.2)

```

```

In [134]: dkt_256_predictions = dkt_256_model.predict([X1_test, X2_test])
          print('DKT Prediction Accuracy with 256 hidden nodes and dropout of 0.01:')
          accuracy_score(np.round(dkt_256_predictions.flatten()), y_test.flatten())

```

DKT Prediction Accuracy with 256 hidden nodes and dropout of 0.01:

```
Out[134]: 0.8621441689623508
```

```

In [137]: print('DKT AUC score:')
          roc_auc_score(y_test.flatten(), dkt_256_predictions.flatten())

```

DKT AUC score:

```
Out[137]: 0.75923135423709032
```

Looks like we were able to get the AUC to increase by 1.7% compared to the original model using 256 hidden nodes and a dropout of 0.1.